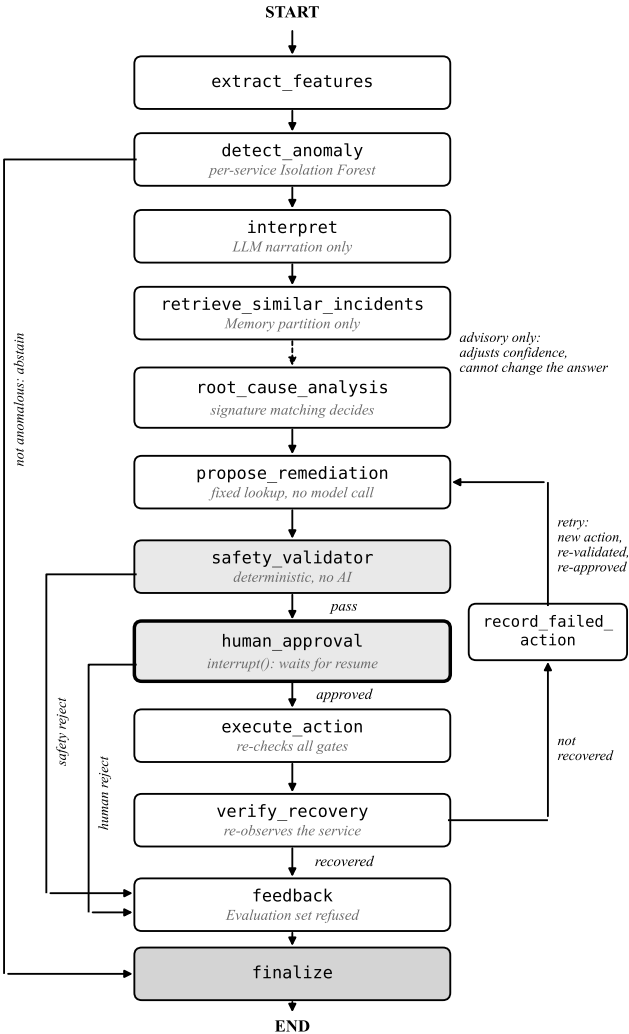
START
extract_features
detect_anomaly
per-service Isolation Forest
interpret
LLM narration only
retrieve_similar_incidents
Memory partition only
advisory only:
adjusts confidence,
cannot change the answer
root_cause_analysis
signature matching decides
propose_remediation
fixed lookup, no model call
retry:
new action,
re-validated,
re-approved
safety_validator
deterministic, no AI
pass
record_failed_
action
human_approval
interrupt(): waits for resume
approved
execute_action
re-checks all gates
not
recovered
verify_recovery
re-observes the service
recovered
feedback
Evaluation set refused
finalize
END
not anomalous: abstain
safety reject
human reject